

Assignment 1 ProblemStatement2 (YARN COMMANDS)

1. Run MapReduce Program and capture the application Id of the job.

```
127 login: march8lab23
march8lab23@127.0.0.1's password:
Last login: Wed Jul 26 12:20:30 2023 from localhost
[march8lab23@ip-10-1-1-204 ~]$ yarn application --list
WARNING: YARN_OPTS has been replaced by HADOOP_OPTS. Using value of YARN_OPTS.
23/07/26 12:46:46 INFO client.RMProxy: Connecting to ResourceManager at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
Total number of applications (application-types: [], states: [SUBMITTED, ACCEPTED, RUNNING] and tags: []):1
Application-Id Application-Name Application-Type User Queue State
Final-State Progress Tracking-URL
application_1685754149182_7570 Spark shell SPARK dhinkumart58edu root.default RUNNING
UNDEFINED 10% http://ip-10-1-1-204.ap-south-1.compute.internal:18088/history/application_1685754149182_7570/1
[march8lab23@ip-10-1-1-204 ~]$ yarn application --list
WARNING: YARN_OPTS has been replaced by HADOOP_OPTS. Using value of YARN_OPTS.
23/07/26 12:47:29 INFO client.RMProxy: Connecting to ResourceManager at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
Total number of applications (application-types: [], states: [SUBMITTED, ACCEPTED, RUNNING] and tags: []):2
Application-Id Application-Name Application-Type User Queue State
Final-State Progress Tracking-URL
application_1685754149182_7606 word count MAPREDUCE march8lab23 root.default ACCEPTED
UNDEFINED 0% N/A
application_1685754149182_7570 Spark shell SPARK dhinkumart58edu root.default RUNNING
UNDEFINED 10% http://ip-10-1-1-204.ap-south-1.compute.internal:18088/history/application_1685754149182_7570/1
[march8lab23@ip-10-1-1-204 ~]$
```

2. Re-run the MapReduce program and kill the application using the yarn command.

```
at java.security.AccessController.doPrivileged(Native Method)
at javax.security.auth.Subject.doAs(Subject.java:422)
at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1875)
at org.apache.hadoop.ipc.Server$Handler.run(Server.java:2680)

at org.apache.hadoop.ipc.Client.getResponse(Client.java:1499)
at org.apache.hadoop.ipc.Client.call(Client.java:1445)
at org.apache.hadoop.ipc.Client.call(Client.java:1355)
at org.apache.hadoop.ipc.ProtobufRpcEngine$Invoker.invoke(ProtobufRpcEngine.java:228)
at org.apache.hadoop.ipc.ProtobufRpcEngine$Invoker.invoke(ProtobufRpcEngine.java:116)
at com.sun.proxy.$Proxy13.getBlockLocations(Unknown Source)
at org.apache.hadoop.hdfs.protocolPB.ClientNamenodeProtocolTranslatorPB.getBlockLocations(ClientNamenodeProtocolTranslatorPB.java:30
1)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:498)
at org.apache.hadoop.io.retry.RetryInvocationHandler.invokeMethod(RetryInvocationHandler.java:422)
at org.apache.hadoop.io.retry.RetryInvocationHandler$Call.invokeMethod(RetryInvocationHandler.java:165)
at org.apache.hadoop.io.retry.RetryInvocationHandler$Call.invoke(RetryInvocationHandler.java:157)
at org.apache.hadoop.io.retry.RetryInvocationHandler$Call.invokeOnce(RetryInvocationHandler.java:95)
at org.apache.hadoop.io.retry.RetryInvocationHandler.invoke(RetryInvocationHandler.java:359)
at com.sun.proxy.$Proxy14.getBlockLocations(Unknown Source)
at org.apache.hadoop.hdfs.DFSClient.callGetBlockLocations(DFSClient.java:853)
... 17 more

23/07/26 12:52:18 INFO mapreduce.Job: map 0% reduce 0%
23/07/26 12:52:18 INFO mapreduce.Job: Job job_1685754149182_7609 failed with state KILLED due to: Application application_1685754149182_7609
was killed by user march8lab23 at 10.1.1.204
23/07/26 12:52:18 INFO mapreduce.Job: Counters: 0
[march8lab23@ip-10-1-1-204 jars]$
```

3. List all the applications which are RUNNING state

```
WARNING: YARN_OPTS has been replaced by HADOOP_OPTS. Using value of YARN_OPTS.
23/07/26 12:50:29 INFO client.RMProxy: Connecting to ResourceManager at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
Exception in thread "main" java.lang.IllegalArgumentException: Invalid ApplicationId prefix: -. The valid ApplicationId should start with pr
efix application
at org.apache.hadoop.yarn.api.records.ApplicationId.fromString(ApplicationId.java:112)
at org.apache.hadoop.yarn.client.cli.ApplicationCLI.killApplication(ApplicationCLI.java:605)
at org.apache.hadoop.yarn.client.cli.ApplicationCLI.killApplication(ApplicationCLI.java:585)
at org.apache.hadoop.yarn.client.cli.ApplicationCLI.run(ApplicationCLI.java:293)
at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:76)
at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:90)
at org.apache.hadoop.yarn.client.cli.ApplicationCLI.main(ApplicationCLI.java:102)
[march8lab23@ip-10-1-1-204 ~]$ yarn application --kill application_1685754149182_7608
WARNING: YARN_OPTS has been replaced by HADOOP_OPTS. Using value of YARN_OPTS.
23/07/26 12:50:41 INFO client.RMProxy: Connecting to ResourceManager at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
Application application_1685754149182_7608 has already finished
[march8lab23@ip-10-1-1-204 ~]$ yarn application --kill application_1685754149182_7609
WARNING: YARN_OPTS has been replaced by HADOOP_OPTS. Using value of YARN_OPTS.
23/07/26 12:52:17 INFO impl.YarnClientImpl: Killed application application_1685754149182_7609
[march8lab23@ip-10-1-1-204 ~]$ yarn application --list --appStates RUNNING
WARNING: YARN_OPTS has been replaced by HADOOP_OPTS. Using value of YARN_OPTS.
23/07/26 12:55:14 INFO client.RMProxy: Connecting to ResourceManager at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
Total number of applications (application-types: [], states: [RUNNING] and tags: []):2
Application-Id Application-Name Application-Type User Queue State
Final-State Progress Tracking-URL
application_1685754149182_7612 word count MAPREDUCE march8lab23 root.default RUNNING
UNDEFINED 5% http://ip-10-1-1-175.ap-south-1.compute.internal:38253
application_1685754149182_7570 Spark shell SPARK dhinkumart58edu root.default RUNNING
UNDEFINED 10% http://ip-10-1-1-204.ap-south-1.compute.internal:18088/history/application_1685754149182_7570/1
[march8lab23@ip-10-1-1-204 ~]$
```

4. View the logs of any of the jobs which are already completed.

```
End of LogType:stdout
*****

End of LogType:prelaunch.err
*****

Container: container_1685754149182_7612_01_000022 on ip-10-1-2-40.ap-south-1.compute.internal_8041
.ogAggregationType: AGGREGATED
*****
.ogType:prelaunch.out
.ogLastModifiedTime:Wed Jul 26 12:56:03 +0000 2023
.ogLength:70
.ogContents:
Setting up env variables
Setting up job resources
Launching container

End of LogType:prelaunch.out
*****

End of LogType:stderr
*****

End of LogType:stdout
*****

[march8lab23@ip-10-1-1-204 ~]$
```